Sparse Matrix Solver

Report

Nourhan Berjawi

I. Mathematical Background and Acknowledgements:

The goal is to solve a system Ax = b.

The iterative methods chosen are the **Jacobi Method** and the **Gauss-Seidel Method**. For both iterative methods to converge, the spectral radius of the matrix must be less than 1. But obtaining the spectral radius can be as time consuming as solving the system with a direct method (such as Gaussian elimination.) To avoid having to do that calculation, and for the solution to converge, it is sufficient that A be strictly diagonally dominant. For the Gauss-Seidel method, it is sufficient for A to be symmetric and positive definite. We know that if a symmetric matrix is strictly row diagonally dominant and has strictly positive diagonal entries, then it is positive definite. Thus, generating a symmetric matrix and ensuring the diagonal entries are strictly dominant is sufficient for both methods to converge.

To achieve this goal, I created a function in main.cpp that, given a size, generates a random sparse symmetric positive matrix, that is also diagonally dominant.

II. Library code breakdown:

The library I have created specializes in solving linear systems involving a sparse matrix, using Compressed Row Storage format. A sparseMatrix<T> object is a templated class, that inherits from Matrix<T>, and that stores a matrix as 3 vectors:
- vector<T> sparse_vals: stores the non-zero matrix vals
- vector<int> cols_ind: stores the corresponding column for every non-zero
- vector<int> rows_ptr: stores the index of where each row begins in sparse_vals and cols_ind

In order to construct a sparseMatrix object, the user must provide the sizes of the matrix as well as a pointer to the array of values. Then, the constructor of sparseMatrix calls the constructor of Matrix, and gets sparse_vals, cols_ind, and rows_ptr from the array of values of the matrix.

Alternatively, the user can initialize an empty sparseMatrix with a specific number of rows and columns, then proceed to fill it later using the set_matrix_vals member function.

The types acceptable to replace T in sparseMatrix<T> are double and int. In other words, the sparseMatrix can be a matrix of doubles or integers. However, the solution returned from solvers will always be a vector<double>.

To solve a system, the user can call the public member functions solve_jacobi and solve_gauss, and given them a vector of integers or doubles.

The Jacobi and Gauss solvers are optimized to only loop over non-zero values, which are stored in the member sparse_vals. They also terminate when the error (sum of absolute differences between $A\bar{x}$ and b) becomes less than 0.5 or when the number of iterations exceeds 5000.

The class sparseMatrix has additional member functions that perform sparse matrix-vector multiplication and getter functions for the members as well as the size. The class has no destructor because vectors were used ensuring garbage collection would be automatically done.

III. Limitations and future improvements:

As mentioned before, the solvers I have implemented converge only for a well-conditioned matrix. They do not manipulate the matrix in any way to overcome the condition problem. In the future, I would like to implement additional solvers that can converge for a larger array of matrices.

IV. main.cpp:

For each of the sizes 10, 50, and 100:
- the main.cpp program generates a random symmetric tridiagonal strictly diagonally dominant matrix *A* and a random vector of doubles *b*
- solves the system using the Jacobi method and displays the solution
- solves the system using the Gauss-Seidel method and displays the solution

To execute main.cpp, download the repository, and in its directory run:

g++ Matrix.cpp sparseMatrix.cpp main.cpp -o main
./main